

# 实验三报告

## 关卡一：openGauss 数据库的编译和安装

### 1. 关卡验证

步骤 1 首先需要对数据库状态进行验证。

```
[omm@opengauss01 openGauss-server]$ gs_ctl status
```

（截图语句和执行结果）

```
[omm@opengauss01 ~]$ gs_ctl status
[2022-12-02 15:36:36.613][227747][][gs_ctl]: gs_ctl status,datadir is /opt/software/openGauss/data
gs_ctl: server is running (PID: 227696)
/opt/software/openGauss/bin/gaussdb "-D" "/opt/software/openGauss/data"
[omm@opengauss01 ~]$
```

步骤 2 对数据库进程进行截图验证，需包含数据库服务器的主机名。

```
[omm@opengauss01 openGauss-server]$ ps -ef|grep omm
```

（截图语句和执行结果）

```
[omm@opengauss01 ~]$ ps -ef|grep omm
root      227476      5364      0 15:25 pts/0    00:00:00 su - omm
omm       227477      227476      0 15:25 pts/0    00:00:00 -bash
omm       227696          1      0 15:33 pts/0    00:00:02 /opt/software/openGauss/bin/gaussdb -D /opt/software/openGauss/data
omm       227796      227477      0 15:38 pts/0    00:00:00 ps -ef
omm       227797      227477      0 15:38 pts/0    00:00:00 grep --color=auto omm
[omm@opengauss01 ~]$
```

## 关卡二：openGauss 数据导入及基本操作

### 1. 关卡验证

步骤 12 登录数据库验证

```
[omm@opengauss01 dbgen]$ gsql -d tpch -p 5432 -r
tpch=# select count(*) from supplier;
```

（截图语句和执行结果）

```
tpch=# select count(*) from supplier;
count
-----
10000
(1 row)

tpch=#
```

## 步骤 21 登录数据库进行验证

```
[omm@opengauss01 ~]$ gsql -d tpch -p 5432 -r
tpch=# \dt
```

(截图语句和执行结果)

```
tpch=# \dt
      List of relations
Schema | Name | Type | Owner | Storage
-----+-----+-----+-----+-----
public | address_dimension | table | omm | {orientation=row,compression=no}
public | customer | table | omm | {orientation=row,compression=no}
public | date_dimension | table | omm | {orientation=row,compression=no}
public | lineitem | table | omm | {orientation=row,compression=no}
public | litemall_orders | table | omm | {orientation=row,compression=no}
public | nation | table | omm | {orientation=row,compression=no}
public | orders | table | omm | {orientation=row,compression=no}
public | part | table | omm | {orientation=row,compression=no}
public | partsupp | table | omm | {orientation=row,compression=no}
public | region | table | omm | {orientation=row,compression=no}
public | supplier | table | omm | {orientation=row,compression=no}
public | user_dimension | table | omm | {orientation=row,compression=no}
(12 rows)

tpch=#
```

## 步骤 22 查询 customer 表的数据

```
tpch=# select * from customer limit 10;
```

(截图语句和执行结果)

```

tpch=# select * from customer limit 10;
 c_custkey |      c_name      |      c_address      | c_nationkey | c_phone | c_acctbal | c_mktsegment |
-----+-----+-----+-----+-----+-----+-----+
1 | Customer#000000001 | IVhzIApeRb ot,c.E | 15 | 25-989-741-2988 | 711.56 | BUILDING | to the even,
regular platelets. regular, ironic epitaphs nag e
2 | Customer#000000002 | XSTf4.NCwDVaWNe6tEgvwfmRchLXak | 13 | 23-768-687-3665 | 121.65 | AUTOMOBILE | l accounts.
blithely ironic theodolites integrate boldly: caref
3 | Customer#000000003 | MG9kdTD2w8Hm | 1 | 11-719-748-3364 | 7498.12 | AUTOMOBILE | deposits ea
t slyly ironic, even instructions. express foxes detect slyly. blithely even accounts abov
4 | Customer#000000004 | XxVSJsLAGtn | 4 | 14-128-198-5944 | 2866.83 | MACHINERY | requests. f
final, regular ideas sleep final accou
5 | Customer#000000005 | KvpyuHCplR8B4WgAiGV6sYpZq7Tj | 3 | 13-758-942-6364 | 794.47 | HOUSEHOLD | n accounts w
ill have to unwind. foxes cajole accor
6 | Customer#000000006 | sKZz8CsnMD7mp4Xd8Yr8vx.LREYKUWAH yVn | 28 | 38-114-968-4951 | 7638.57 | AUTOMOBILE | tions. even
deposits boost according to the slyly bold packages. final accounts cajole requests. furious
7 | Customer#000000007 | TcGe5gaZNgVePxUSKRrvXBfkasDTea | 18 | 28-198-982-9759 | 9561.95 | AUTOMOBILE | ainst the ir
onic, express theodolites. express, even pinto beans among the exp
8 | Customer#000000008 | I8B18bB8AymmC. 0PrRYBCPiYGJ8xc8PmWhl5 | 17 | 27-147-574-9335 | 6819.74 | BUILDING | among the sl
yly regular theodolites kindle blithely courts. carefully even theodolites haggle slyly along the ide
9 | Customer#000000009 | xKiAFTJUsCuxfeleNgefumTrjS | 8 | 18-338-986-3675 | 8324.87 | FURNITURE | r theodolite
s according to the requests wake thinly excuses: pending requests haggle furiousl
10 | Customer#000000010 | 6LrEav6KR6PLVcgl2Arl Q3rqzLzcT1 v2 | 5 | 15-741-346-9878 | 2753.54 | HOUSEHOLD | es regular d
eposits haggle. fur
(18 rows)

tpch=#

```

## 2. 思考题

数据初始化中出现了 TPC-H，这是什么？

答：TPC-H 是用于生成测试数据的测试包。

## 关卡三：openGauss 的 AI4DB 特性应用

### 1. 关卡验证

#### (1) 使用 X-Tuner 进行参数优化

步骤 2 在原来 CloudShell 连接窗口中查看 queries01.log。

```
[omm@opengauss01 ~]$ tail -10 /opt/software/tpch-kit/dbgen/queries/queries01.log
```

（截图执行语句和结果）

```

[omm@opengauss01 ~]$ tail -10 /opt/software/tpch-kit/dbgen/queries/queries01.log
13 |      | 888 | 6737713.99
17 |      | 861 | 6460573.72
18 |      | 964 | 7236687.40
23 |      | 892 | 6701457.95
29 |      | 948 | 7158866.63
30 |      | 909 | 6808436.13
31 |      | 922 | 6806670.18
(7 rows)

total time: 1181031 ms
[omm@opengauss01 ~]$

```

步骤 3 切换至 root 用户，执行 X-Tuner 进行参数建议优化

```
[omm@opengauss01 ~]$ exit
[root@opengauss01 xtuner]# gs_xtuner recommend --db-name tpch --db-user omm --port 5432
--host 127.0.0.1 --host-user omm
```

（截图执行语句和结果）

```
***** Recommended Knob Settings *****
```

name	recommend	min	max	restart
default_statistics_target	1000	100	1000	False
effective_cache_size	21602334	186756	21602334	False
effective_io_concurrency	200	150	250	False
enable_mergejoin	off	0	1	False
enable_nestloop	off	0	1	False
max_connections	370	50	741	True
max_prepared_transactions	370	50	741	True
max_process_memory	28803112	22402420	28803112	True
random_page_cost	1.0	1.0	2.0	False
shared_buffers	186756	186760	214772	True
wal_buffers	5836	2048	5836	True

```
[root@opengauss01 xtuner]#
```

步骤 6 获取参数值

```
[omm@opengauss01 ~]$ cd /opt/software/openGauss/data
[omm@opengauss01 data]$ cat postgresql.conf|grep -E
'shared_buffers|max_connections|effective_cache_size|effective_io_concurrency|wal_buffers|random_page_cost|default_statistics_target'
```

（截图执行语句和结果）

```
[omm@opengauss01 data]$ cat postgresql.conf|grep -E 'shared_buffers|max_connections|effective_cache_size|effective_io_concurrency|wal_buffers|random_page_cost|default_statistics_target'
# (change requires restart)
max_connections = 370
# Note: Increasing max_connections costs ~400 bytes of shared memory per
shared_buffers = 186756 # min 128kB
bulk_write_ring_size = 200 # for bulkload, max shared_buffers
standby_shared_buffers_fraction = 0.3 # control shared buffers use in standby, 0.1-1.0
effective_io_concurrency = 200 # 1-1000; 0 disables prefetching
wal_buffers = 5836 # min 32kB
random_page_cost = 1 # same scale as above
effective_cache_size = 21602334
default_statistics_target = 1000 # range 1-10000
# max_locks_per_transaction * (max_connections + max_prepared_transactions)
[omm@opengauss01 data]$
```

步骤 7 再次执行步骤 2，对比优化前的执行时间。

（截图执行语句和结果）

```
[omm@opengauss01 ~]$ tail -10 /opt/software/tpch-kit/dbgen/queries/queries01.log
13 | 888 | 6737713.99
17 | 861 | 6460573.72
18 | 964 | 7236687.40
23 | 892 | 6701457.95
29 | 948 | 7158866.63
30 | 909 | 6808436.13
31 | 922 | 6806670.18
(7 rows)

total time: 1158289 ms
```

步骤 8 【附加题】有兴趣的同学可以尝试并截图记录于此。

(截图执行语句和结果)

```
[omm@opengauss01 ~]$ tail -10 /opt/software/tpch-kit/dbgen/queries/queries01.log
13      |      888 | 6737713.99
17      |      861 | 6460573.72
18      |      964 | 7236687.40
23      |      892 | 6701457.95
29      |      948 | 7158866.63
30      |      909 | 6808436.13
31      |      922 | 6806670.18
(7 rows)

total time: 321706 ms
[omm@opengauss01 ~]$
```

## (2) Index-advisor: 索引推荐

步骤 4 使用 explain, 对该 SQL 加以分析

```
tpch=# EXPLAIN
SELECT ad.province AS province, SUM(o.actual_price) AS GMV
FROM litemall_orders o,
     address_dimension ad,
     date_dimension dd
WHERE o.address_key = ad.address_key
     AND o.add_date = dd.date_key
     AND dd.year = 2020
     AND dd.month = 3
GROUP BY ad.province
ORDER BY SUM(o.actual_price) DESC;
```

(截图执行语句和结果)

```
tpch=# tpch=# EXPLAIN
SELECT ad.province AS province, SUM(o.actual_price) AS GMV
FROM litemall_orders o,
     address_dimension ad,
     date_dimension dd
WHERE o.address_key = ad.address_key
     AND o.add_date = dd.date_key
     AND dd.year = 2020
     AND dd.month = 3
GROUP BY ad.province
ORDER BY SUM(o.actual_price) DESC;

----- QUERY PLAN -----

Sort (cost=4593.80..4593.88 rows=31 width=47)
  Sort Key: (sum(o.actual_price)) DESC
  -> HashAggregate (cost=4592.72..4593.03 rows=31 width=47)
    Group By Key: ad.province
    -> Hash Join (cost=4354.43..4585.97 rows=1351 width=15)
      Hash Cond: (ad.address_key = o.address_key)
      -> Seq Scan on address_dimension ad (cost=0.00..188.02 rows=8002 width=14)
      -> Hash (cost=4337.54..4337.54 rows=1351 width=9)
        -> Hash Join (cost=1031.78..4337.54 rows=1351 width=9)
          Hash Cond: (o.add_date = dd.date_key)
          -> Seq Scan on litemall_orders o (cost=0.00..3041.00 rows=100000 width=13)
          -> Hash (cost=1031.76..1031.76 rows=2 width=4)
            -> Seq Scan on date_dimension dd (cost=0.00..1031.76 rows=2 width=4)
              Filter: ((year = 2020) AND ((month)::bigint = 3))

(14 rows)

tpch=#
```



步骤 9 使用 explain, 对该 SQL 加以分析

```
tpch=# EXPLAIN
SELECT ad.province AS province, SUM(o.actual_price) AS GMV
FROM litemall_orders o,
     address_dimension ad,
     date_dimension dd
WHERE o.address_key = ad.address_key
     AND o.add_date = dd.date_key
     AND dd.year = 2020
     AND dd.month = 3
GROUP BY ad.province
ORDER BY SUM(o.actual_price) DESC;
```

(截图执行语句和结果)

```
tpch=# EXPLAIN
tpch=# SELECT ad.province AS province, SUM(o.actual_price) AS GMV
tpch=# FROM litemall_orders o,
tpch=# address_dimension ad,
tpch=# date_dimension dd
tpch=# WHERE o.address_key = ad.address_key
tpch=# AND o.add_date = dd.date_key
tpch=# AND dd.year = 2020
tpch=# AND dd.month = 3
tpch=# GROUP BY ad.province
tpch=# ORDER BY SUM(o.actual_price) DESC;

----- QUERY PLAN -----

Sort (cost=3579.58..3579.65 rows=31 width=47)
  Sort Key: (sum(o.actual_price)) DESC
  -> HashAggregate (cost=3578.58..3578.61 rows=31 width=47)
    Group By Key: ad.province
    -> Hash Join (cost=3340.21..3571.74 rows=1351 width=15)
      Hash Cond: (ad.address_key = o.address_key)
      -> Seq Scan on address_dimension ad (cost=0.00..188.02 rows=8002 width=14)
      -> Hash (cost=3323.32..3323.32 rows=1351 width=9)
        -> Hash Join (cost=17.56..3323.32 rows=1351 width=9)
          Hash Cond: (o.add_date = dd.date_key)
          -> Seq Scan on litemall_orders o (cost=0.00..3841.00 rows=100000 width=13)
          -> Hash (cost=17.53..17.53 rows=2 width=4)
            -> Index Scan using <16583>btree_date_dimension_year on date_dimension dd (cost=0.00..17.53 rows=2 width=4)
              Index Cond: (year = 2020)
              Filter: ((month)::bigint = 3)

(15 rows)

tpch=#
```

步骤 11 【附加题】有兴趣的同学可以尝试并截图记录于此。

(截图执行语句和结果)

```
[omm@opengauss01 ~]$ tail -10 /opt/software/tpch-kit/dbgen/queries/queries02.log
13      |      888 | 6737713.99
17      |      861 | 6460573.72
18      |      964 | 7236687.40
23      |      892 | 6701457.95
29      |      948 | 7158866.63
30      |      909 | 6808436.13
31      |      922 | 6806670.18
(7 rows)

total time: 320980 ms
[omm@opengauss01 ~]$
```

## 关卡四 【附加题】： openGauss 的 DB4AI 特性应用

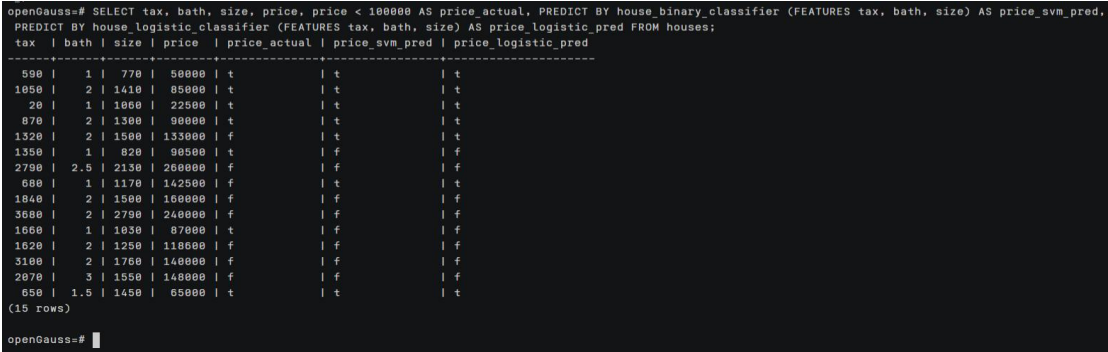
\*本关卡为附加题，有兴趣的同学可以尝试实验并记录于此。

# 1. 关卡验证

步骤 10 利用训练好的逻辑回归模型预测数据，并与 SVM 算法进行比较，将执行结果截图。

```
openGauss=# SELECT tax, bath, size, price, price < 100000 AS price_actual, PREDICT BY
house_binary_classifier (FEATURES tax, bath, size) AS price_svm_pred, PREDICT BY
house_logistic_classifier (FEATURES tax, bath, size) AS price_logistic_pred FROM houses;
```

（截图执行语句和结果）



## 清理工作：资源释放

# 1. 关卡验证

步骤 3 查看到列表中已没有资源时，表示弹性云服务器已删除。

（截图执行语句和结果）

